

# Dive into Deep Learning for NLP

## 0. Introduction

Haibin Lin and Leonard

[gluon-nlp.mxnet.io](http://gluon-nlp.mxnet.io)

8:30-9:00	Continental Breakfast
9:00-9:45	Introduction and Setup
9:45-10:30	Neural Networks 101
10:30-10:45	Break
10:45-11:15	Machine Learning Basics
11:15-11:45	Context-free Representations for Language
11:45-12:15	Convolutional Neural Networks
12:15-13:15	Lunch Break
13:15-14:00	Recurrent Neural Networks
14:00-14:45	Attention Mechanism and Transformer
14:45-15:00	Coffee Break
15:00-16:15	Contextual Representations for Language
16:15-17:00	Language Generation

# Deep Learning

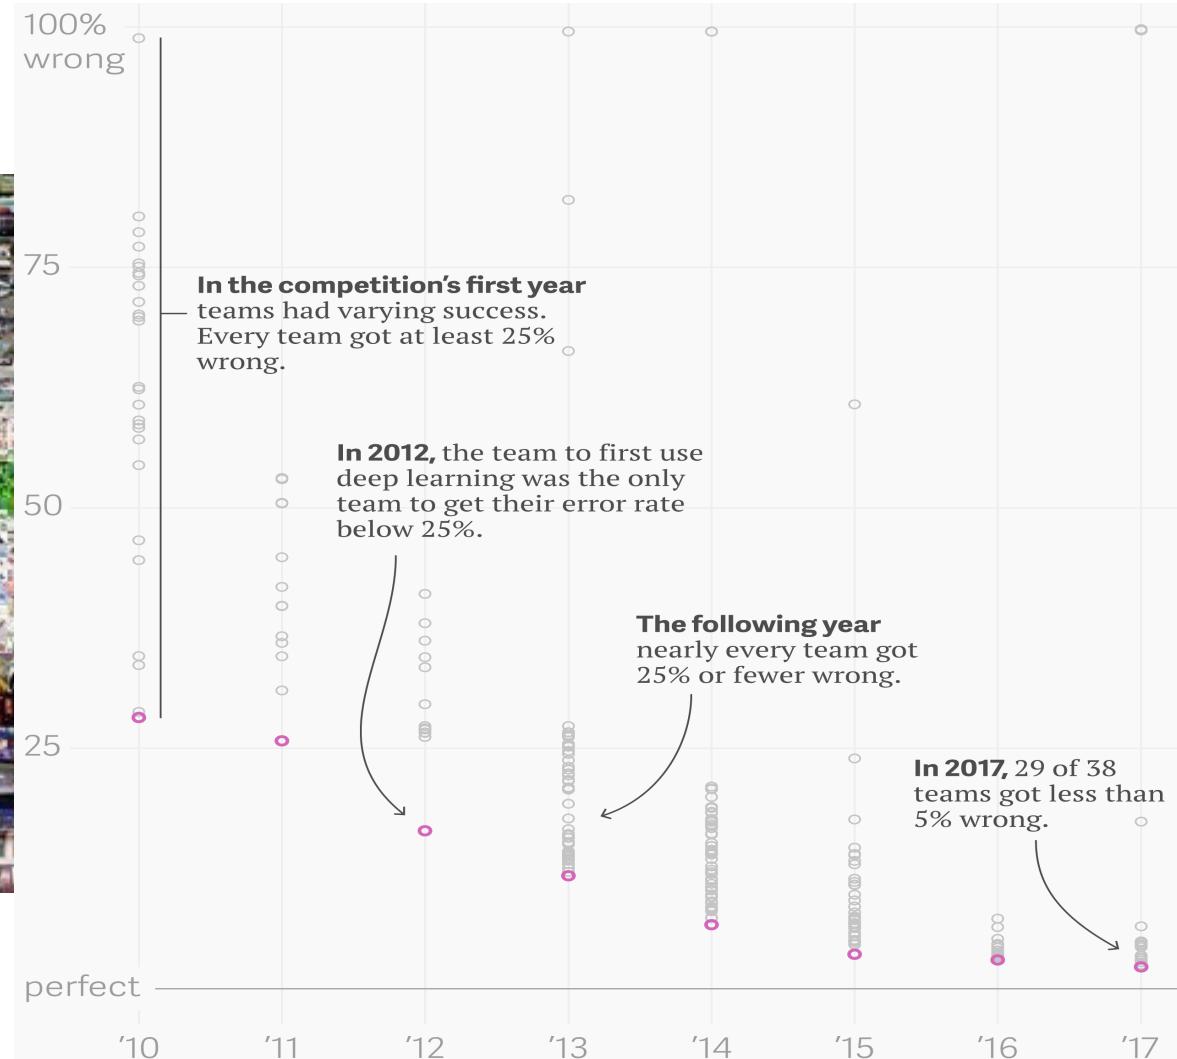


# Classify Images



<http://www.image-net.org/>

# Classify Images



Yanofsky, Quartz

<https://qz.com/1034972/the-data-that-changed-the-direction-of-ai-research-and-possibly-the-world/>

# Detect and Segment Objects



# Style transfer

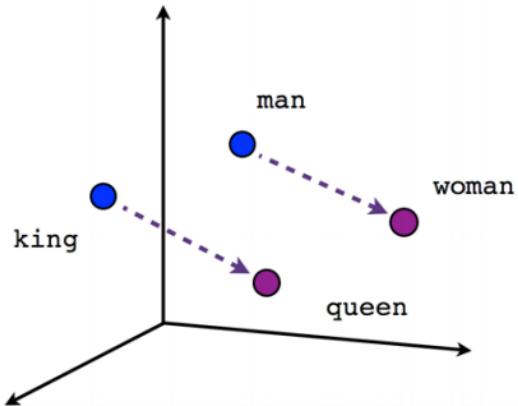


<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>

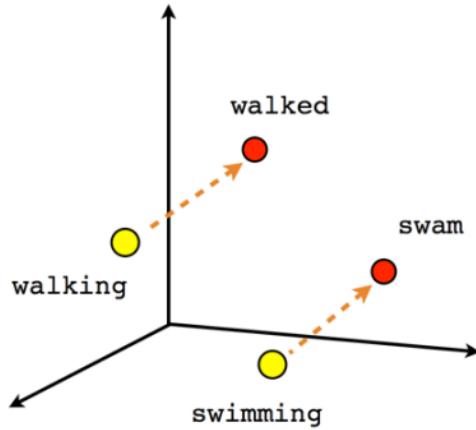
# Synthesize Images



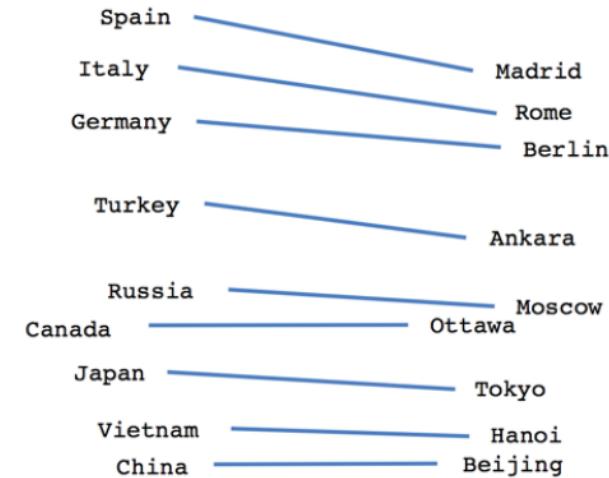
# Analogies



Male-Female



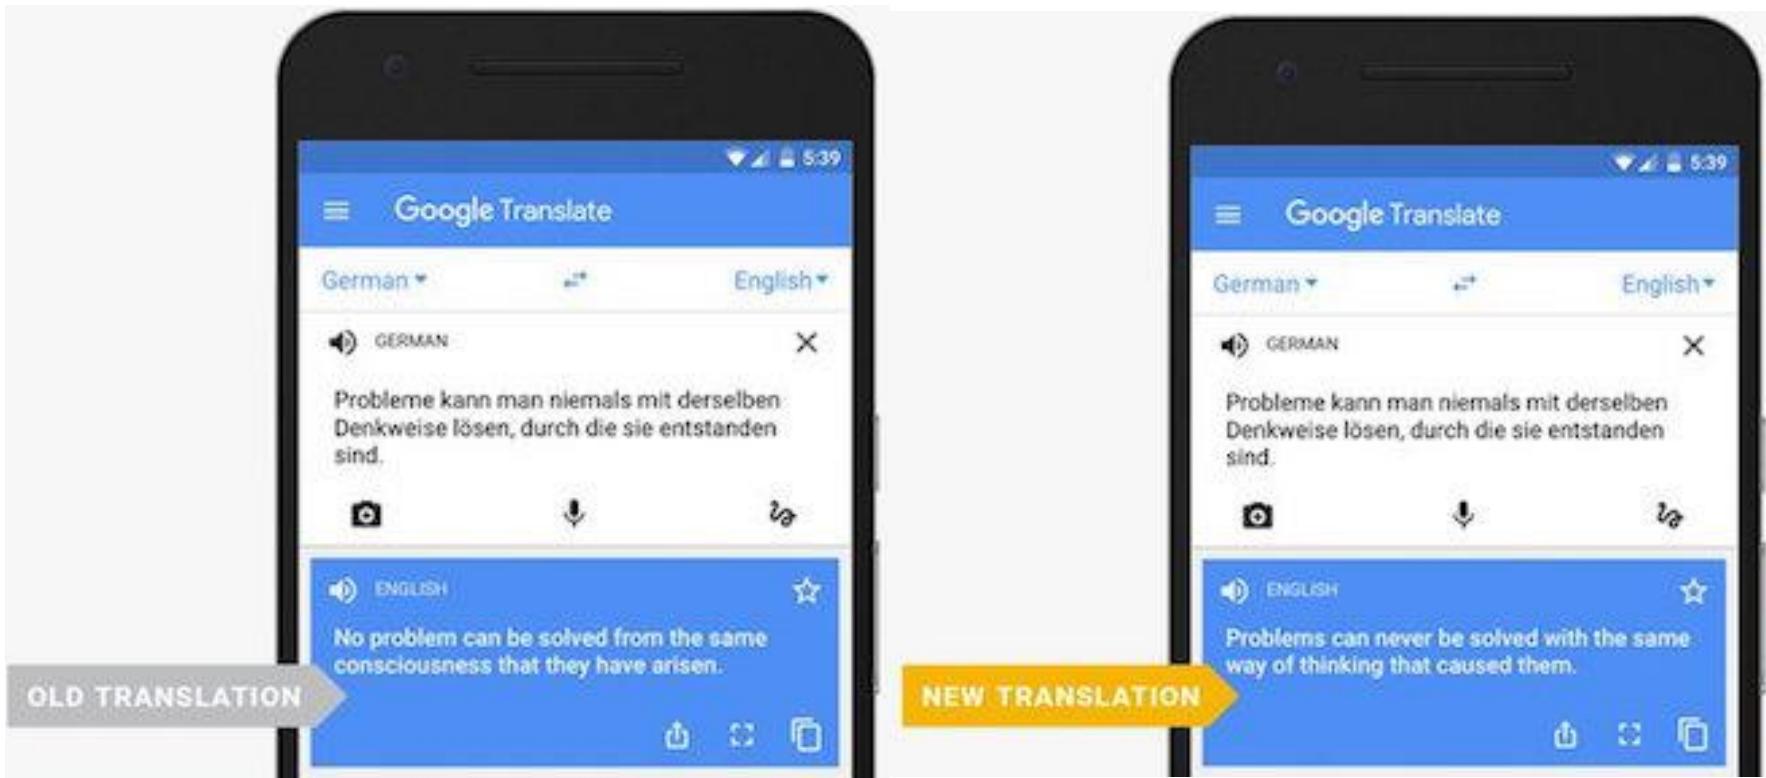
Verb tense



Country-Capital

<https://www.tensorflow.org/tutorials/word2vec>

# Machine Translation



<https://www.pcmag.com/news/349610/google-expands-neural-networks-for-language-translation>

# Text synthesis

**Content:** Two dogs play by a tree.

**Style:** **happily, love**



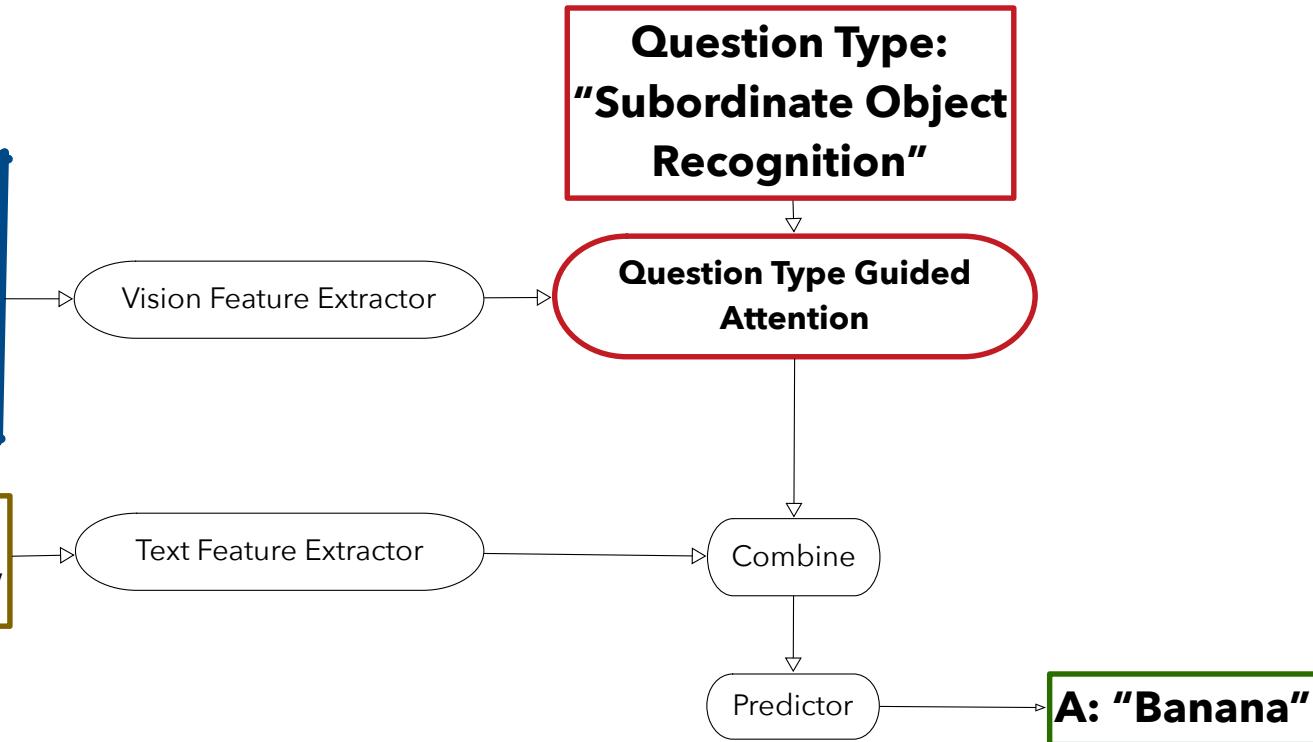
Two dogs **in love** play **happily** by a tree.

Li et al, NACCL, 2018

# Question answering



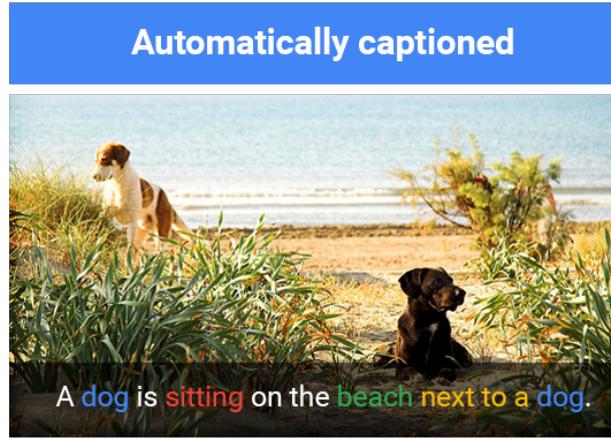
**Q: "What's her  
mustache made of?"**



Shi et al, 2018, Arxiv

# Image captioning

Human captions from the training set



Shallue et al, 2016

<https://ai.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>



# Software

# Programming for deep learning used to be...

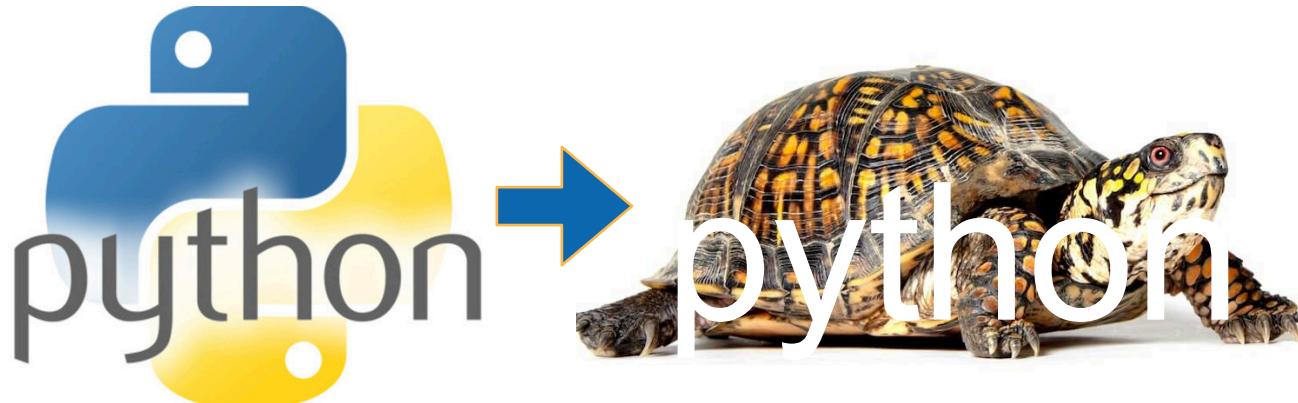
```
void computeLogregSoftmaxGrad(NVMatrix& labels, NVMatrix& probs, NVMatrix& target, bool add, float coeff) {  
    int numCases = probs.getLeadingDim();  
    int numOut = probs.getFollowingDim();  
    assert(labels.getNumElements() == numCases);  
    assert(probs.isContiguous());  
    assert(target.isContiguous());  
    assert(labels.isContiguous());  
    assert(probs.isTrans());  
  
    dim3 threads(LOGREG_GRAD_THREADS_X, LOGREG_GRAD_THREADS_Y);  
    dim3 blocks(DIVUP(numCases, LOGREG_GRAD_THREADS_X), DIVUP(numOut, LOGREG_GRAD_THREADS_Y));  
    if (!add) {  
        target.resize(probs);  
        kLogregSoftmaxGrad<false><<<blocks, threads>>>(probs.getDevData(), labels.getDevData(), target.getDevData(),  
                                                               numCases, numOut, coeff);  
    } else {  
        kLogregSoftmaxGrad<true><<<blocks, threads>>>(probs.getDevData(), labels.getDevData(), target.getDevData(),  
                                                               numCases, numOut, coeff);  
    }  
  
    getLastCudaError("computeLogregSoftmaxGrad: Kernel execution failed");  
}
```

Krizhevsky et al, code from cuda-convnet (AlexNet)

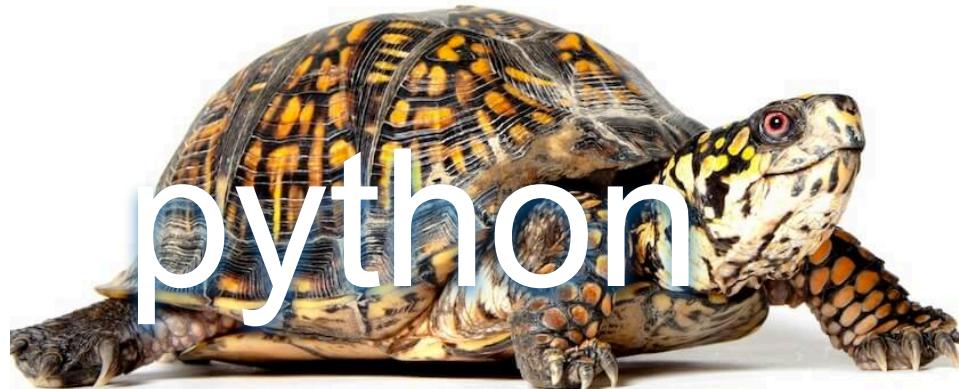


# Programming for deep learning today

```
loss_function = gluon.loss.SoftmaxCrossEntropyLoss()  
loss = loss_function(output)  
loss.backward()
```



# Programming for deep learning today



PYTORCH



theano

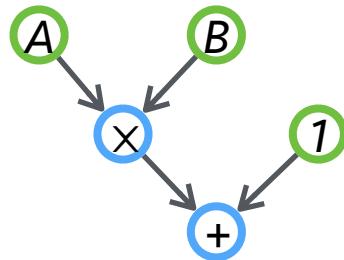
```
output = model(data)
```



[crazy GPU stuff]

aws

# Declarative (Symbolic) Programs



```
A = Variable('A')  
B = Variable('B')  
C = B * A  
D = C + 1  
f = compile(D)  
d = f(A=np.ones(10),  
      B=np.ones(10)*2)
```

- **Advantages:**

- Opportunities for optimization
- Easy to serialize models
- More portable across languages

- **Disadvantages:**

- Hard to debug
- Unsuitable for dynamic graphs
- Can't use native code

C can share memory with D

# Imperative Programs



```
import numpy as np  
a = np.ones(10)  
b = np.ones(10) * 2  
c = b * a  
print c  
d = c + 1
```

Easy to debug,  
easy to code

- **Advantages**

- Straightforward and flexible
- Use language natively (loops, control flow, debugging)

- **Disadvantages**

- Hard to optimize  
(Fixed with JIT compiler!)

# MXNet Gluon

Q: Why GLUON ?

A:



# NDArray



# N-dimensional Array Examples

- N-dimensional array, short for ndarray, is the main data structure for machine learning and neural networks

0-d (scalar)



1.0

A class label

1-d (vector)



[1.0, 2.7, 3.4]

A feature vector

2-d (matrix)



[[1.0, 2.7, 3.4]  
 [5.0, 0.2, 4.6]  
 [4.3, 8.5, 0.2]]

A batch-by-feature matrix

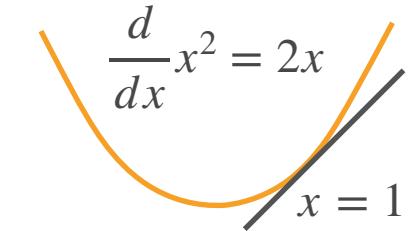
# Review Scalar Derivative

$y$	$a$	$x^n$	$\exp(x)$	$\log(x)$	$\sin(x)$
<hr/>					
$\frac{dy}{dx}$	0	$nx^{n-1}$	$\exp(x)$	$\frac{1}{x}$	$\cos(x)$

( $a$  is not a function of  $x$ )

$y$	$u + v$	$uv$	$y = f(u), u = g(x)$
<hr/>			
$\frac{dy}{dx}$	$\frac{du}{dx} + \frac{dv}{dx}$	$\frac{du}{dx}v + \frac{dv}{dx}u$	$\frac{dy}{du} \frac{du}{dx}$

Derivative is the slope of the tangent line



The slope of the tangent line is 2

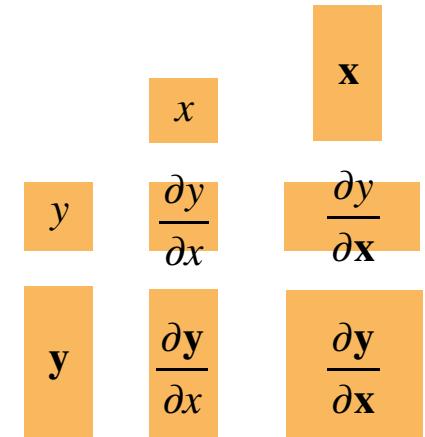
# Gradients

- Generalize derivatives into vectors

Vector		
Scalar		
Scalar	$x$	$\mathbf{x}$
Scalar	$y$	$\frac{\partial y}{\partial x}$
Vector	$\mathbf{y}$	$\frac{\partial \mathbf{y}}{\partial x}$

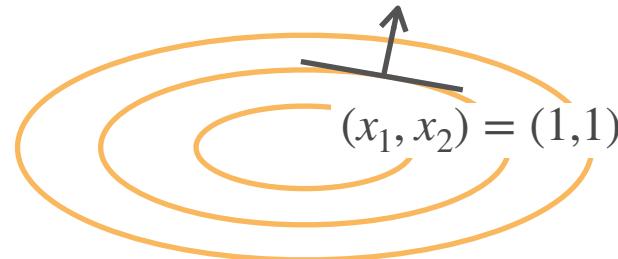
$\partial y / \partial \mathbf{x}$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \frac{\partial y}{\partial \mathbf{x}} = \left[ \frac{\partial y}{\partial x_1}, \frac{\partial y}{\partial x_2}, \dots, \frac{\partial y}{\partial x_n} \right]$$



Direction (2, 4), perpendicular  
to the contour lines

$$\frac{\partial}{\partial \mathbf{x}} x_1^2 + 2x_2^2 = [2x_1, 4x_2]$$



# Generalize to Vectors

- Chain rule for scalars:

$$y = f(u), \quad u = g(x) \quad \frac{\partial y}{\partial x} = \frac{\partial y}{\partial u} \frac{\partial u}{\partial x}$$

- Generalize to vectors straightforwardly

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathbf{y}}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$$

(1,n) (1,) (1,n)

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathbf{y}}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$$

(1,n) (1,k) (k, n)

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathbf{y}}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$$

(m, n) (m, k) (k, n)

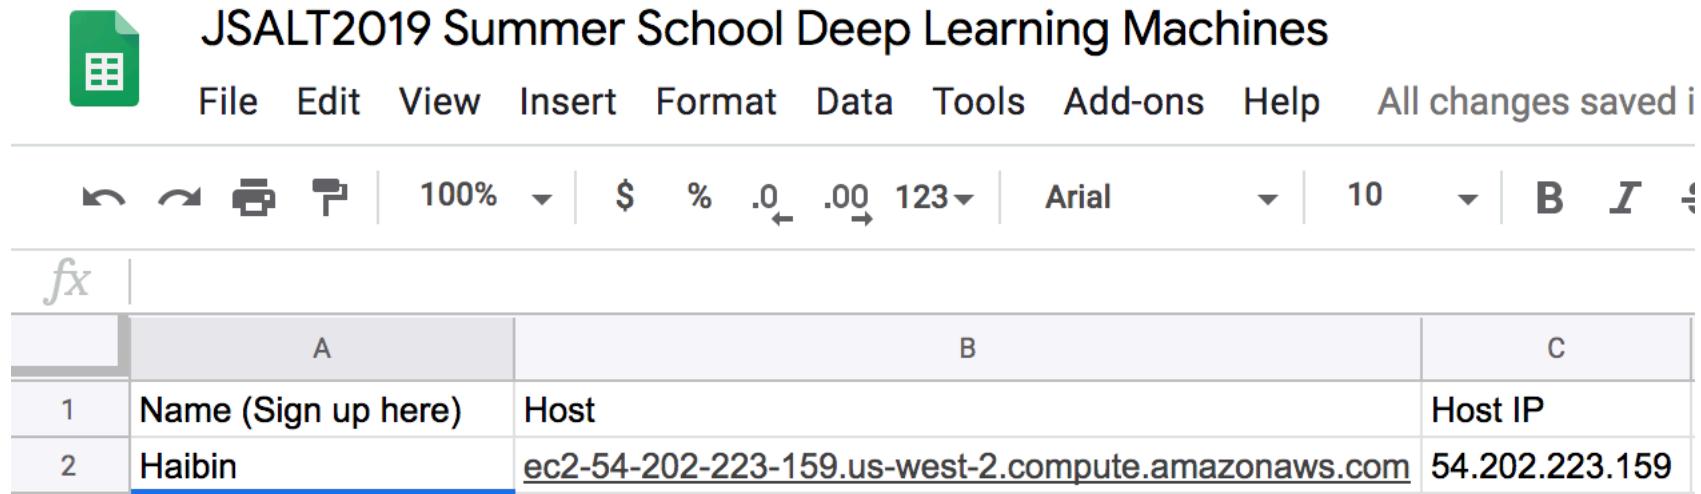


# Environment Setup

- **Setup Instructions**
  - <https://tinyurl.com/JSALT2019>
- **Tutorial Website**
  - <https://jsalt19.mxnet.io>
- **Related Resources**
  - GluonNLP <http://gluon-nlp.mxnet.io/>
  - GluonCV <http://gluon-cv.mxnet.io/>
  - Deep Graph Library <http://dgl.ai/>
  - Dive into deep learning book <http://d2l.ai/>
  - MXNet forum <https://discuss.mxnet.io/>

# Grab a Machine

- Fill in your name in <https://tinyurl.com/jsalt-gpu>



	A	B	C
1	Name (Sign up here)	Host	Host IP
2	Haibin	<u>ec2-54-202-223-159.us-west-2.compute.amazonaws.com</u>	54.202.223.159

# Connect to the Instance

- download the key from [secret link here]
- ssh to the machine

```
chmod 400 jsalt19.pem;
```

```
ssh -i jsalt19.pem ubuntu@your_ip -L 8888:localhost:8888
```

```
→ Downloads chmod 400 jsalt19.pem; ssh -i jsalt19.pem ubuntu@54.202.223.159 -L 8888:localhost:8888
```



```
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.4.0-1084-aws x86_64v)
```

```
Apache MXNet Website: https://beta.mxnet.io
```

```
D2L.ai: http://d2l.ai
```

```
GluonNLP: https://gluon-nlp.mxnet.io
```

```
GluonCV: https://gluon-cv.mxnet.io
```

```
Enjoy deep learning!
```

```
gluon-nlp.mxnet.io
```



# Prepare Environment

- Create a tmux session

```
tmux new -s gluonnlp
```

- Create a conda environment

```
conda env update --prune mxnet/jsalt19; conda activate jsalt19
```

```
Successfully uninstalled gluonnlp-0.7.0.dev0
Successfully installed gluonnlp-0.7.0.dev0 mxnet-cu92mkl-1.5.0b20190613
```

```
#
# To activate this environment, use
#
#     $ conda activate jsalt19
#
# To deactivate an active environment, use
#
#     $ conda deactivate
```

```
ubuntu@ip-172-31-4-238:~$ conda activate jsalt19
(jsalt19) ubuntu@ip-172-31-4-238:~$ █
```



# Clone and Launch Notebooks

git clone <https://github.com/eric-haibin-lin/JSALT19-GluonNLP>;  
jupyter notebook

```
(jsalt19) ubuntu@ip-172-31-4-238:~$ jupyter notebook
[I 02:12:29.882 NotebookApp] Using EnvironmentKernelSpecManager...
[I 02:12:29.882 NotebookApp] Started periodic updates of the kernel list (every 3 minutes).
[I 02:12:30.923 NotebookApp] [jupyter_nbextensions_configurator] enabled 0.4.1
[I 02:12:30.923 NotebookApp] Serving notebooks from local directory: /home/ubuntu
[I 02:12:30.923 NotebookApp] The Jupyter Notebook is running at:
[I 02:12:30.923 NotebookApp] http://localhost:8888/?token=a0f8f026225a211eed7ee86f5d2f6dd926c82aee4dbda8e1
[I 02:12:30.923 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[W 02:12:30.927 NotebookApp] No web browser found: could not locate runnable browser.
[C 02:12:30.928 NotebookApp]
```

To access the notebook, open this file in a browser:

file:///run/user/1000/jupyter/nbserver-8650-open.html

Or copy and paste one of these URLs:

<http://localhost:8888/?token=a0f8f026225a211eed7ee86f5d2f6dd926c82aee4dbda8e1>

```
[I 02:12:30.928 NotebookApp] Starting initial scan of virtual environments...
```

